

# Chapter 2: Operating-System Structures

---



# Outline

---

## Outline

- Operating System Services
- User Operating System Interface
- System Calls
- System Programs
- Operating System Structure

## Objectives

- To describe the services an operating system provides to users, processes, and other systems
- To discuss the various ways of structuring an operating system

# Operating System Services

---

- **For user**
  - User interface
  - Program execution
  - I/O operations
  - File-system manipulation
  - Communication process communication
  - Error detection and handling
- **For System:** efficiency and sharing
  - Resource allocation
  - Accounting
  - Protection and security

# Operating System Services

---

- Operating systems provide an environment for execution of programs and services to programs and users
- One set of operating-system services provides functions that are helpful to the user:
  - **User interface** - Almost all operating systems have a user interface (**UI**).
    - Varies between **Command-Line (CLI)**, **Graphics User Interface (GUI)**, **Batch**
  - **Program execution** - The system must be able to load a program into memory and to run that program, end execution, either normally or abnormally (indicating error)
  - **I/O operations** - A running program may require I/O, which may involve a file or an I/O device

# Operating System Services (Cont.)

---

- One set of operating-system services provides functions that are helpful to the user (Cont.):
  - **File-system manipulation** - The file system is of particular interest. Programs need to read and write files and directories, create and delete them, search them, list file information, permission management.
  - **Communications** – Processes may exchange information, on the same computer or between computers over a network
    - Communications may be via shared memory or through message passing (packets moved by the OS)
  - **Error detection** – OS needs to be constantly aware of possible errors
    - May occur in the CPU and memory hardware, in I/O devices, in user program
    - For each type of error, OS should take the appropriate action to ensure correct and consistent computing
    - Debugging facilities can greatly enhance the user's and programmer's abilities to efficiently use the system

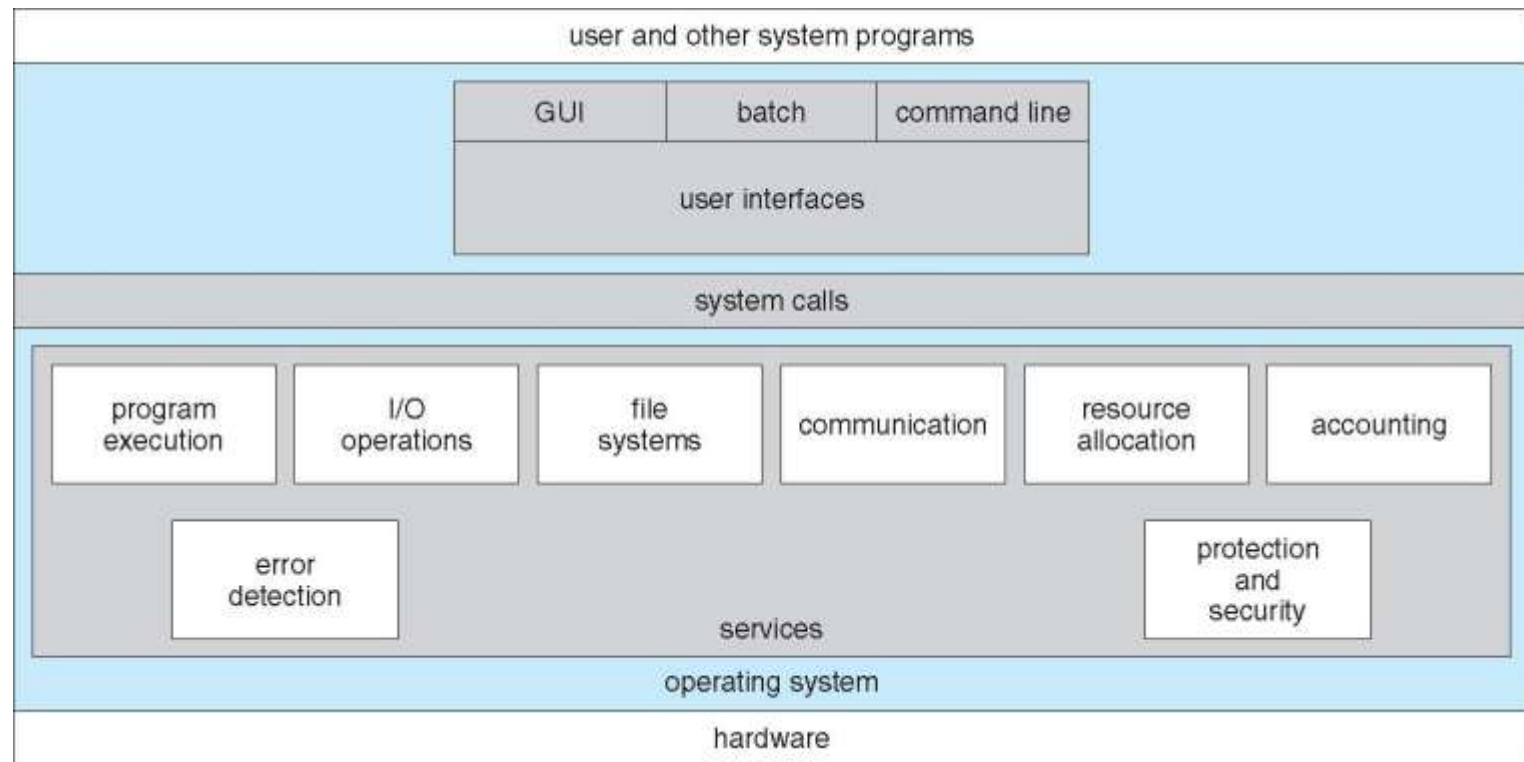
# Operating System Services (Cont.)

---

- Another set of OS functions exists for ensuring the efficient operation of the system itself via resource sharing
  - **Resource allocation** - When multiple users or multiple jobs running concurrently, resources must be allocated to each of them
    - Many types of resources - CPU cycles, main memory, file storage, I/O devices.
  - **Accounting** - To keep track of which users use how much and what kinds of computer resources
  - **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other
    - **Protection** involves ensuring that all access to system resources is controlled
    - **Security** of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts

# A View of Operating System Services


---



# [User - Operating System] Interface - CLI

---

- CLI: Command Line Interface (CLI) or **command interpreter (shell)**
  - in kernel or as a system program,
  - Many flavors
  - fetches a command from user and executes it
    - Command may be **built-in**,
    - Command may be **another program**
- GUI: User-friendly **desktop** interface
  - **Icons** represent files, programs, actions, etc.



Many operating systems now include both CLI and GUI interfaces  
Linux: command shells available (CLI); KDE as GUI



# Bourne Shell Command Interpreter

```
Default
New Info Close Execute Bookmarks

PBG-Mac-Pro:~ pbg$ w
15:24 up 56 mins, 2 users, load averages: 1.51 1.53 1.65
USER      TTY      FROM          LOGIN@   IDLE   WHAT
pbg       console -              14:34    50    -
pbg       s000    -              15:05    -    w

PBG-Mac-Pro:~ pbg$ iostat 5

            disk0             disk1             disk10             cpu             load average
      KB/t tps  MB/s      KB/t tps  MB/s      KB/t tps  MB/s  us sy id   1m   5m   15m
      33.75 343 11.30      64.31 14   0.88      39.67  0   0.02 11 5 84  1.51 1.53 1.65
        5.27 320  1.65         0.00  0   0.00         0.00  0   0.00  4 2 94  1.39 1.51 1.65
        4.28 329  1.37         0.00  0   0.00         0.00  0   0.00  5 3 92  1.44 1.51 1.65

^C
PBG-Mac-Pro:~ pbg$ ls
Applications                Music                        WebEx
Applications (Parallels)    Pando Packages             config.log
Desktop                      Pictures                     getsmartdata.txt
Documents                   Public                      imp
Downloads                   Sites                       log
Dropbox                     Thumbs.db                   panda-dist
Library                     Virtual Machines             prob.txt
Movies                      Volumes                     scripts

PBG-Mac-Pro:~ pbg$ pwd
/Users/pbg

PBG-Mac-Pro:~ pbg$ ping 192.168.1.1
PING 192.168.1.1 (192.168.1.1): 56 data bytes
64 bytes from 192.168.1.1: icmp_seq=0 ttl=64 time=2.257 ms
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=1.262 ms
^C
--- 192.168.1.1 ping statistics ---
2 packets transmitted, 2 packets received, 0.0% packet loss
round-trip min/avg/max/stddev = 1.262/1.760/2.257/0.498 ms
PBG-Mac-Pro:~ pbg$
```

# User Operating System Interface - GUI

---

- User-friendly **desktop** metaphor interface
  - Usually mouse, keyboard, and monitor
  - **Icons** represent files, programs, actions, etc
  - Various mouse buttons over objects in the interface cause various actions (provide information, options, execute function, open directory (known as a **folder**))
  - Invented at Xerox PARC
- Many systems now include both CLI and GUI interfaces
  - Microsoft Windows is GUI with CLI “command” shell
  - Apple Mac OS X is “Aqua” GUI interface with UNIX kernel underneath and shells available
  - Unix and Linux have CLI with optional GUI interfaces (CDE, KDE, GNOME)

The collage features several overlapping screenshots of the Mac OS X graphical user interface:

- Top Left:** A screenshot of the **Finder** application window showing a folder named "My Documents". The sidebar lists various locations like Home, Applications, Desktop, Library, etc.
- Top Center:** A screenshot of a **Finder** window displaying a file named "Rg-206.tiff". The right pane shows detailed metadata for the TIFF file, including creation date, size, and format.
- Bottom Left:** A screenshot of the **iPhoto** application window, showing a library of photos and a grid view.
- Bottom Center:** A screenshot of the **Apple Computer Inc.** website as it appeared in the early 2000s, featuring the company logo and contact information.
- Bottom Right:** A screenshot of the **Dictionary** application window, showing the definition for the term "operating system".

The background of the collage is a blurred image of a Mac OS X desktop with various icons and a clock widget.

# Touchscreen Interfaces

- ☐ Touchscreen devices require new interfaces
  - ☐ Mouse not possible or not desired
  - ☐ Actions and selection based on gestures
  - ☐ Virtual keyboard for text entry
- ☐ Voice commands.



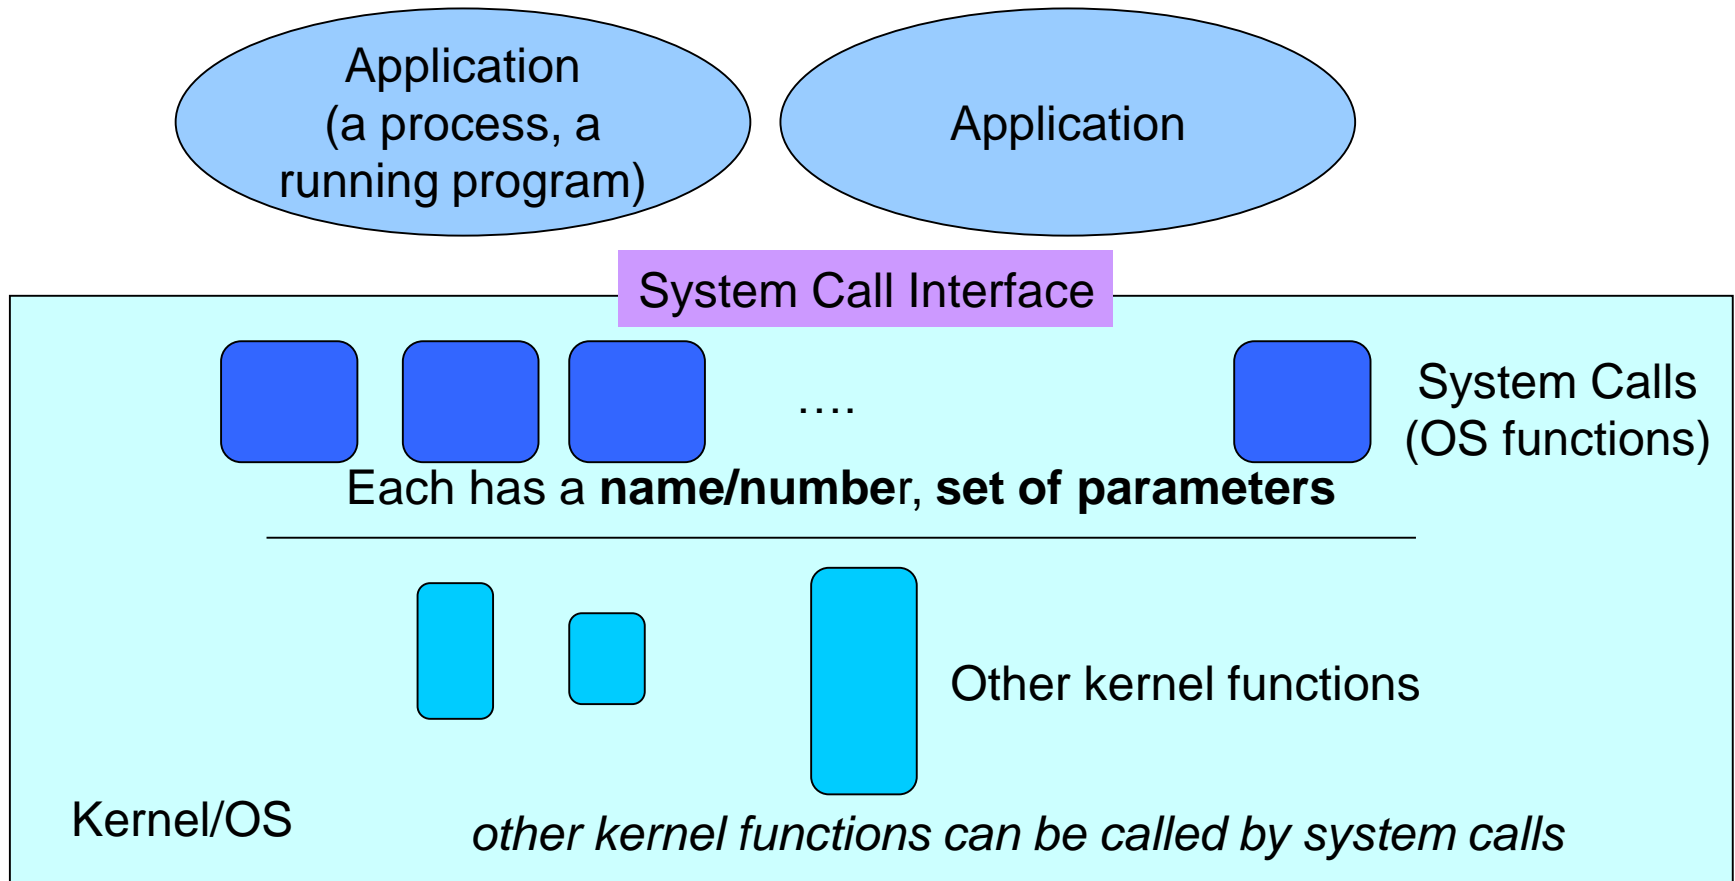
# System Calls

---

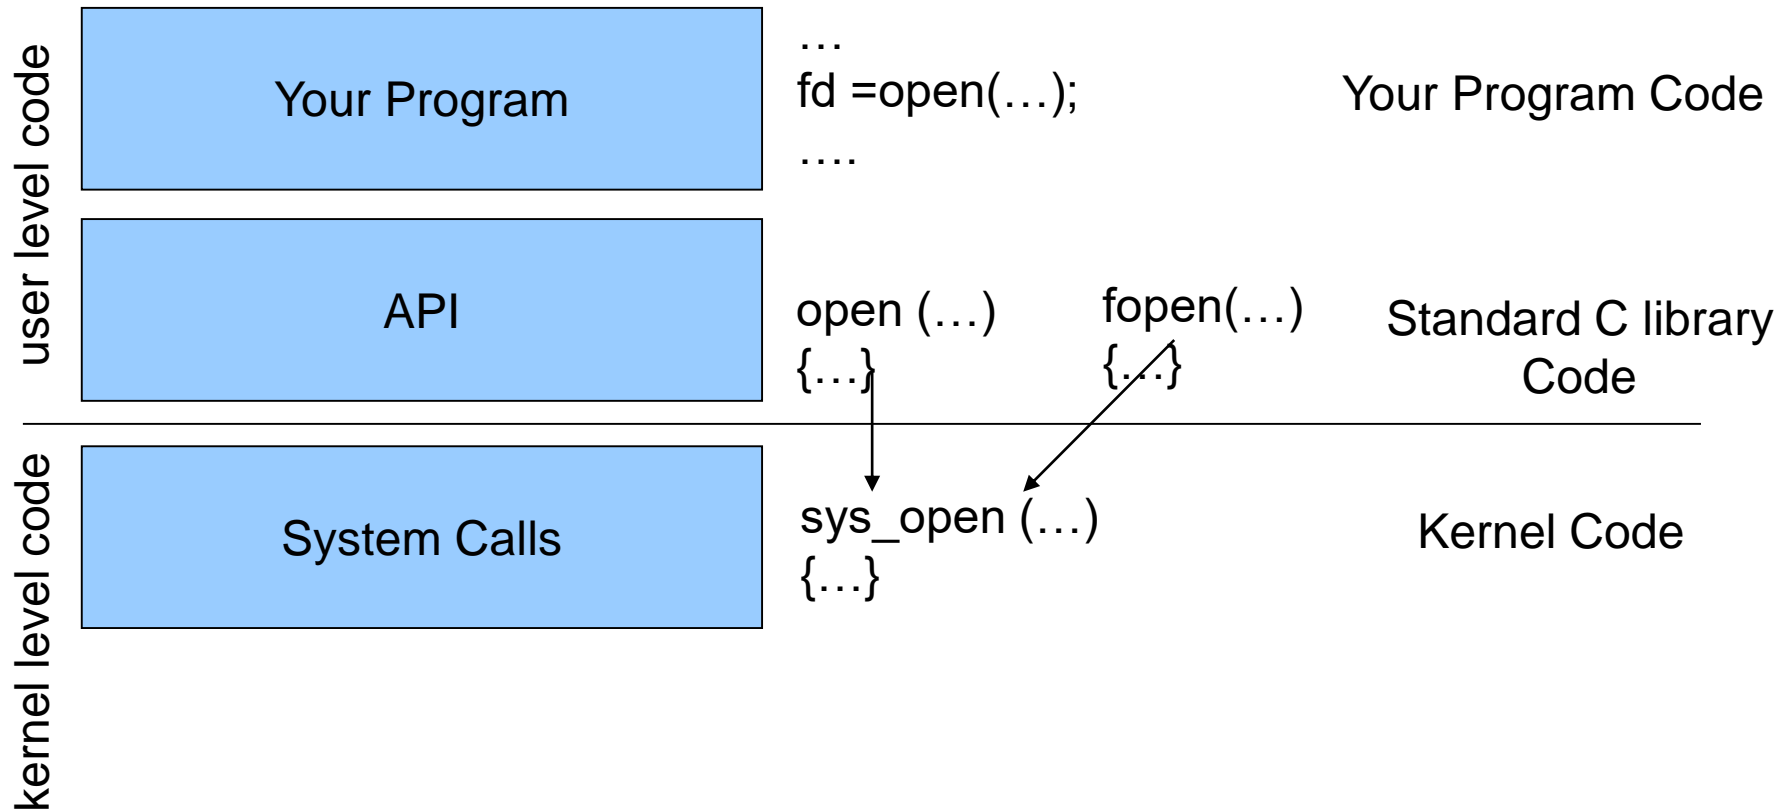
- Programming interface to the services provided by the OS
- Typically written in a high-level language (C or C++)
- Mostly accessed by programs via a high-level **Application Programming Interface (API)** rather than direct system call use
- Three most common APIs are
- Win32 API for Windows, POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X), and Java API for the Java virtual machine (JVM)

Note that the system-call names used throughout this text are generic

# System Calls



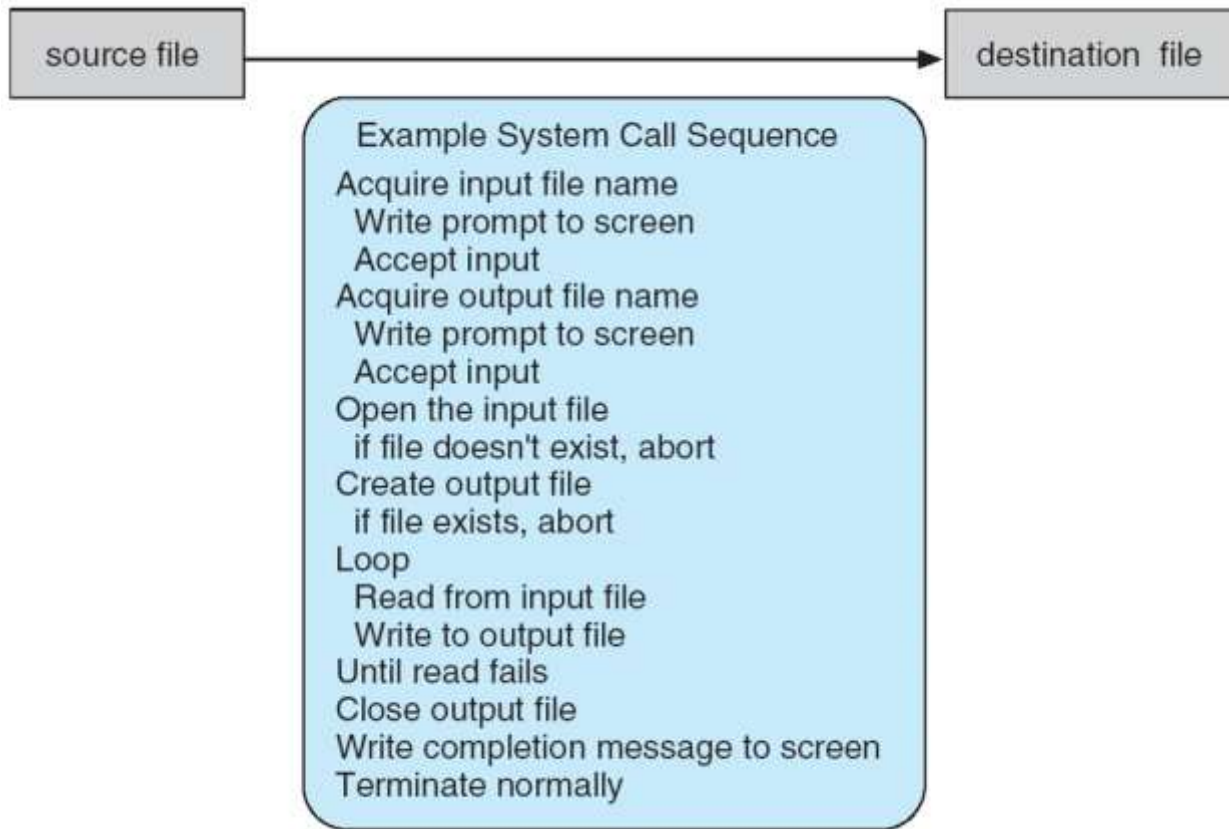
# System Calls



# Example of System Calls

---

- System call sequence to copy the contents of one file to another file






# Example of Standard API

---


- Consider the ReadFile() function in the Win32 API — a function for reading from a file



- A description of the parameters passed to ReadFile()
  - HANDLE file—the file to be read
  - LPVOID buffer—a buffer where the data will be read into and written from
  - DWORD bytesToRead—the number of bytes to be read into the buffer
  - LPDWORD bytesRead—the number of bytes read during the last read
  - LPOVERLAPPED ovl—indicates if overlapped I/O is being used

# System Call Implementation and Calling

- Typically,
  - a **number** associated with each system call
  - Number used as an index to a table: System Call **table**
  - Table keeps **addresses** of system calls (routines)
  - System call runs and returns
- Caller does not know system call implementation
  - Just knows interface



# Linux system calls

---


around 300-400 system calls

Number	Generic Name of System Call	Name of Function in Kernel
1	exit	sys_exit
2	fork	sys_fork
3	read	sys_read
4	write	sys_write
5	open	sys_Open
6	close	sys_Close
...		
39	mkdir	sys_Mkdir
...		

# Standard C Library Example

---

- C program invoking printf() library call, which calls write() system call



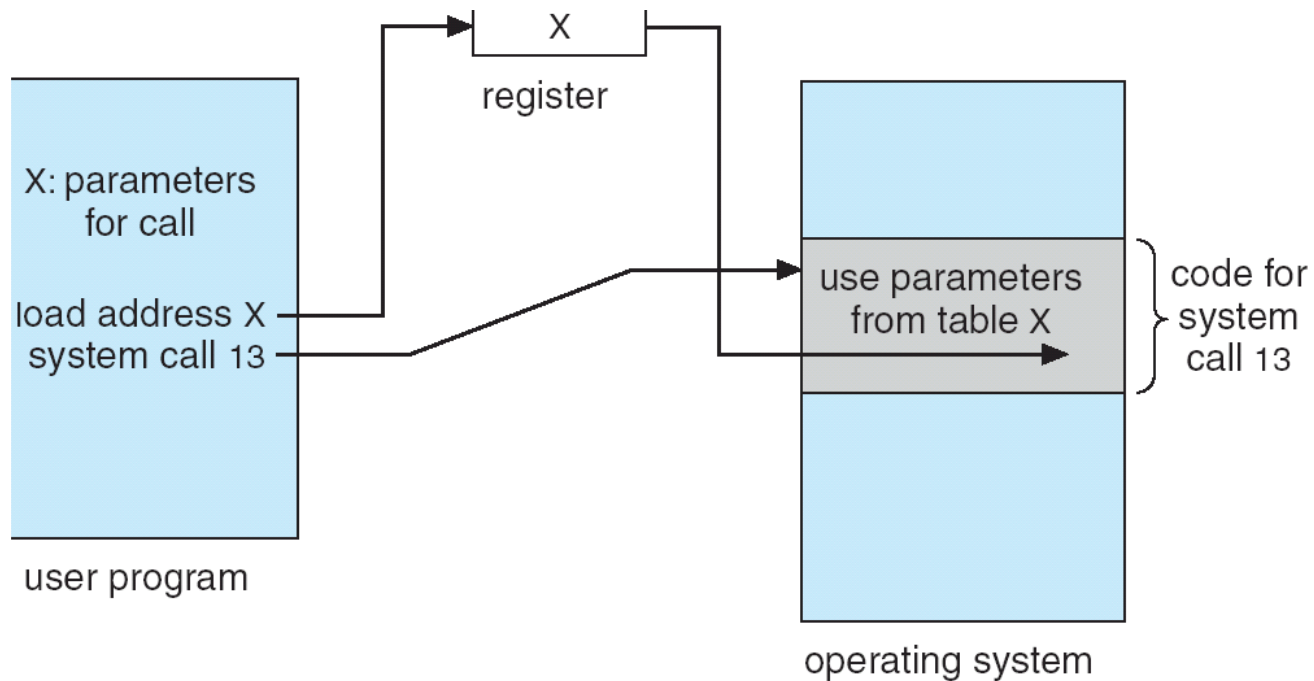
# System Call Parameter Passing

---

- Often, more information is required than the identity of the desired system call
  - Exact type and amount of information vary according to OS and call
- Three general methods used to pass parameters to the OS
  - 1) Simplest: pass the parameters in *registers*
    - In some cases, may be more parameters than registers
  - 2) Parameters stored in a *block*, or *table*, in memory, and address of block passed as a parameter in a register
  - 3) Parameters placed, or *pushed*, onto the *stack* by the program and *popped* off the stack by the operating system

Last two methods do not limit the number or length of parameters being passed

# Parameter Passing via Table



# Types of System Calls

---

- Process control
- File management
- Device management
- Information maintenance
- Communications
- Protection

# System Call types Examples of Windows and Unix System Calls

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()



# System Programs

---

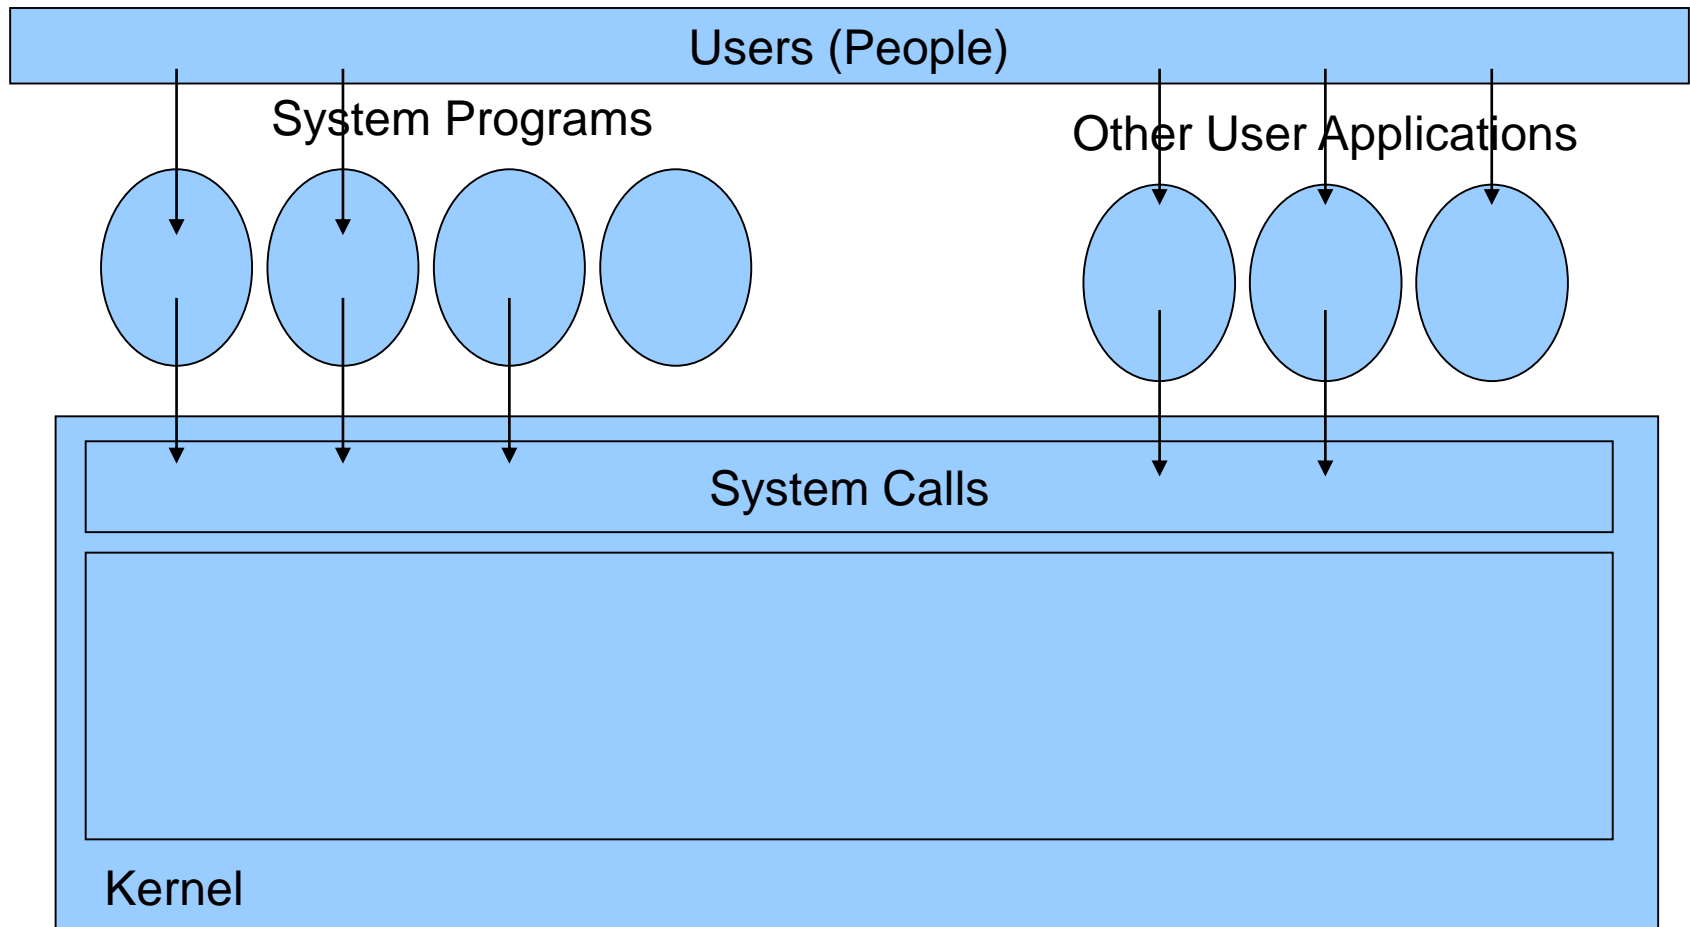
- System programs provide a convenient environment for program development and execution. They can be divided into:
  - File manipulation (create, delete, copy, rename, print, list, ...)
  - Status information (date, time, amount of available memory, disk space, who is logged on, ...)
  - File modification (text editors, grep, ...)
  - Programming language support (compiler, debuggers, ...)
  - Program loading and execution (loaders, linkers)
  - Communications (ftp, browsers, ssh, ...)
  - Other System Utilities/Applications may come with OS CD (games, math solvers, plotting tools, database systems, spreadsheets, word processors, ...)

# System Programs

---

- Most users' view of the operation system is defined by system programs, not the actual system calls
- System programs provide a convenient environment for program development and execution
- Some of them are simply user interfaces to system calls; others are considerably more complex
  - create file: simple system program that can just call “create” system call or something similar
  - compiler: complex system program

# System Programs



From OS' s view: system+user programs are all applications

---

# Structuring OS (Kernel)

# OS Structure

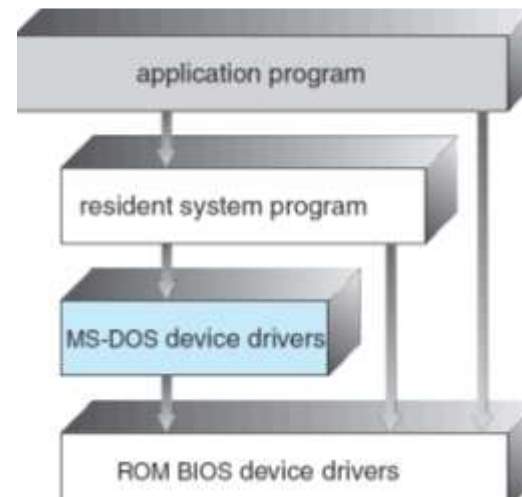
---

- Simple Structure (MSDOS)
- Layered Approach
- Microkernel Approach
- Modules Approach

# Simple Structure

---


- MS-DOS – written to provide the most functionality in the least space
  - Not divided into modules
  - Although MS-DOS has some structure, its interfaces and levels of functionality are not well separated



# Layered Approach

---

- The operating system is divided into a number of layers (levels), each built on top of lower layers. The bottom layer (layer 0), is the hardware; the highest (layer N) is the user interface.
- With modularity, layers are selected such that each uses functions (operations) and services of only lower-level layers



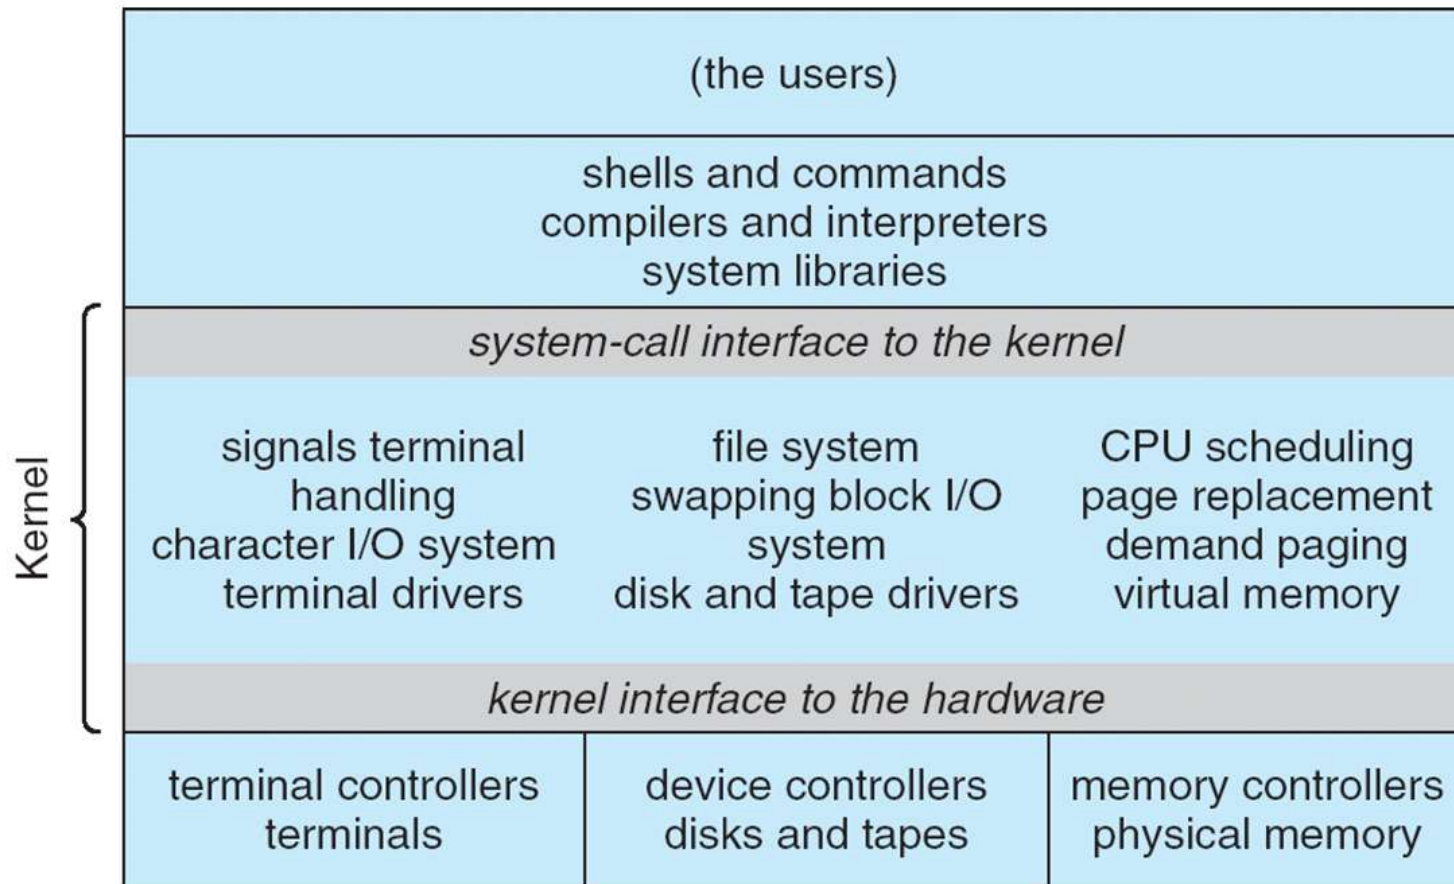
# Unix

---

- UNIX – limited by hardware functionality, the original UNIX operating system had limited structuring. The UNIX OS consists of two separable parts
  - Systems programs
  - The kernel
    - Consists of everything below the system-call interface and above the physical hardware
    - Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level



# Traditional UNIX System Structure

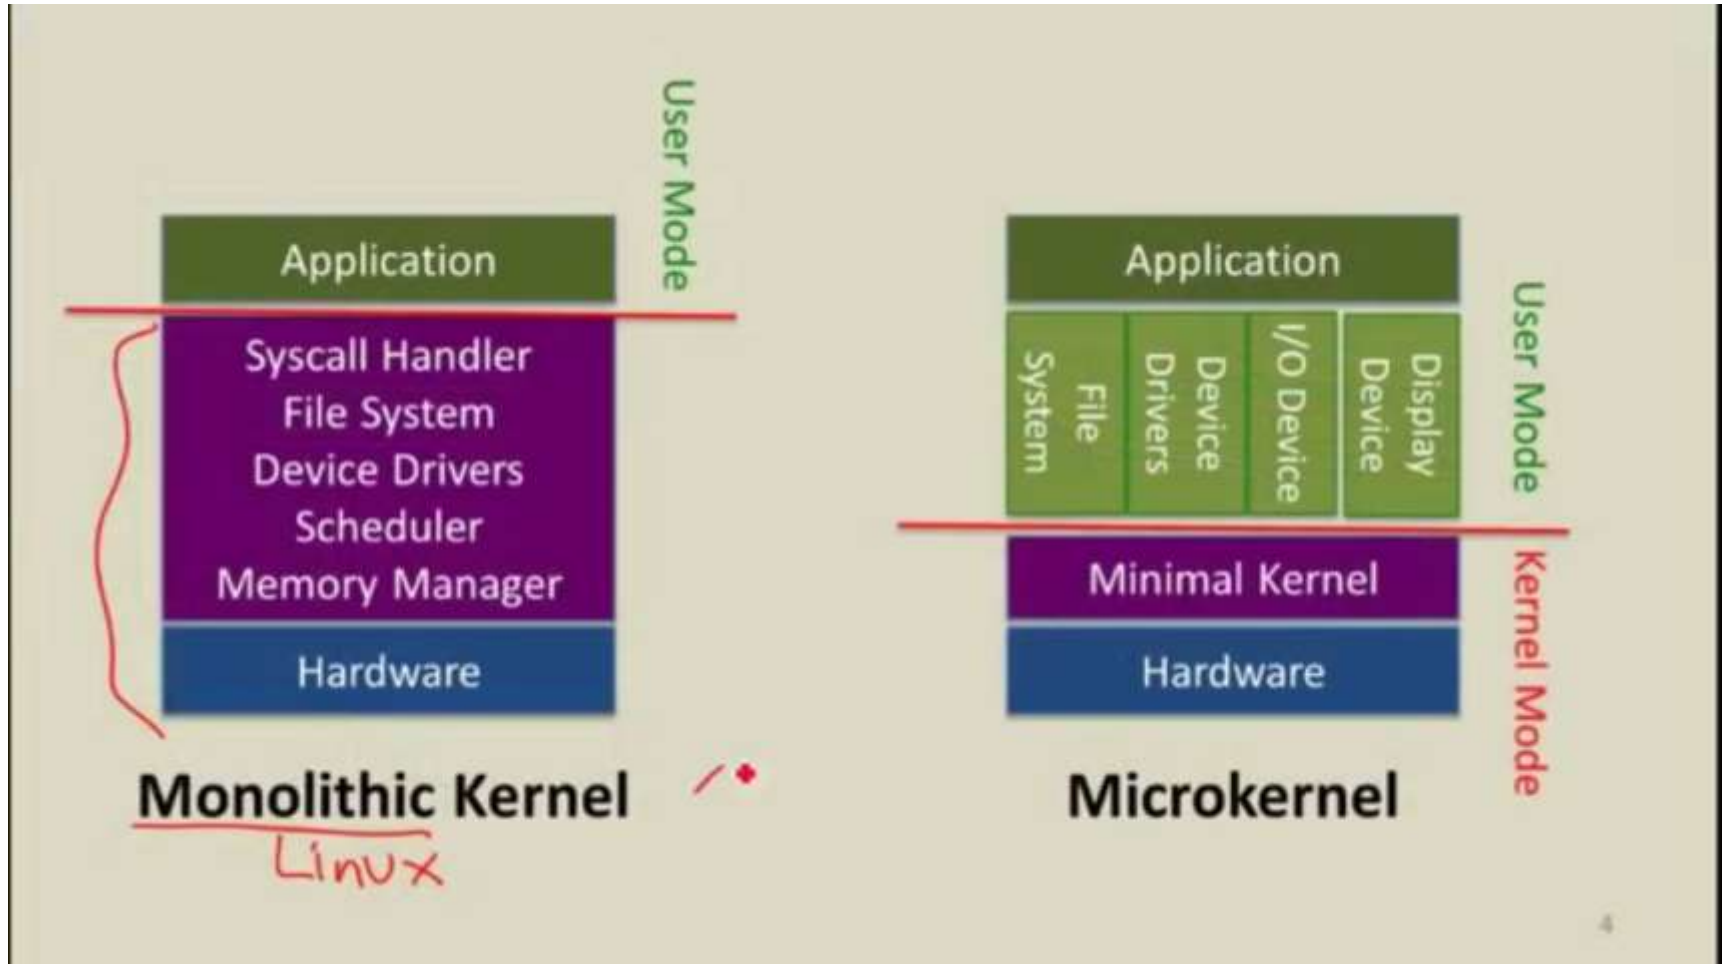


# Microkernel System Structure

---

- Moves as much from the kernel into “*user*” space
- Communication takes place between user modules using message passing
- Benefits:
  - Easier to extend a microkernel
  - Easier to port the operating system to new architectures
  - More reliable (less code is running in kernel mode)
  - More secure
- Detriments:
  - Performance overhead of user space to kernel space communication

# Microkernel Vs Monolithic Kernel



# Microkernel Vs Monolithic Kernel

## 1. MICROKERNEL VS MONOLITHIC SYSTEM

~1992

Most older operating systems are monolithic, that is, the whole operating system is a single a.out file that runs in 'kernel mode.' This binary contains the process management, memory management, file system and the rest. Examples of such systems are UNIX, MS-DOS, VMS, MVS, OS/360, MULTICS, and many more.

The alternative is a microkernel-based system, in which most of the OS runs as separate processes, mostly outside the kernel. They communicate by message passing. The kernel's job is to handle the message passing, interrupt handling, low-level process management, and possibly the I/O. Examples of this design are the RC4000, Amoeba, Chorus, Mach, and the not-yet-released Windows/NT.

While I could go into a long story here about the relative merits of the two designs, suffice it to say that among the people who actually design operating systems, the debate is essentially over. **Microkernels have won.**

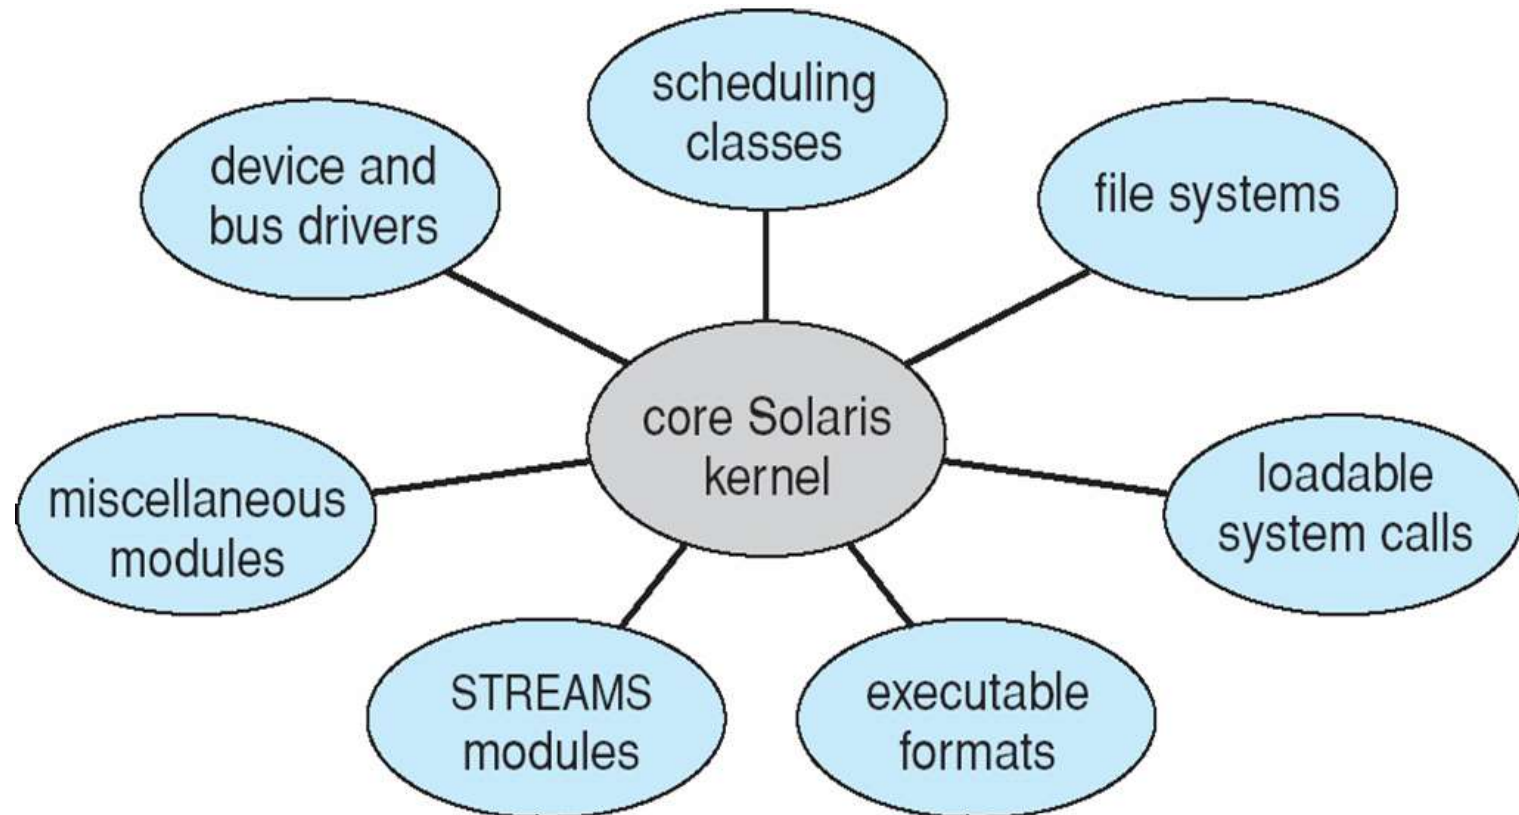
# Modules

---

- Most modern operating systems implement kernel modules
  - Uses object-oriented approach
  - Each core component is separate
  - Each talks to the others over known interfaces
  - Each is **loadable** as needed within the kernel
- Overall, similar to layers but more flexible
- Linux supports modules

# Solaris Modular Approach

---



---

# **Additional Study Material**

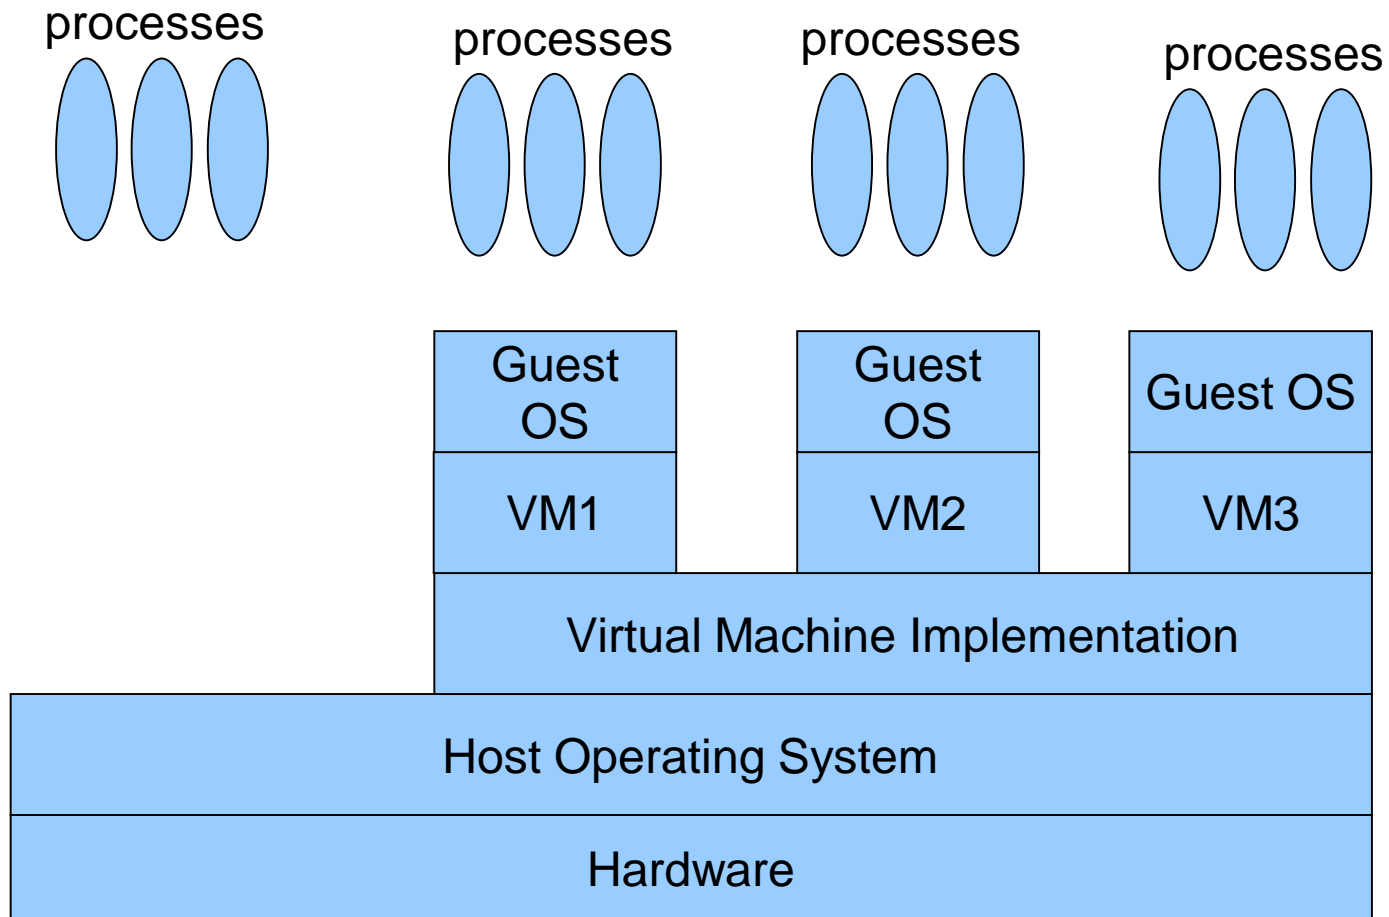
# Virtual Machines

---

- Hardware is abstracted into several different execution environments
  - Virtual machines
- Each virtual machine provides an interface that is identical to the bare hardware
- A *guest* process/kernel can run on top of a virtual machine.
  - We can run several operating systems on the same *host*.
  - Each virtual machine will run another operating system.



# Virtual Machines



# Examples

---

- VMware
  - Abstracts Intel X86 hardware
- Java virtual machine
  - Specification of an abstract computer
- .NET Framework

# Operating System Debugging

---

- Failure analysis
  - Log files
  - Core dump
  - Crash dump
- Performance tuning
  - Monitor system performance
    - Add code to kernel
    - Use system tools like “top”
- DTrace
  - Facility to dynamically adding probes to a running system (both to processes and to the kernel)
  - Probes can be queries using D programming language to obtain info

# Operating System Generation

---

- Configure the kernel
- Compile the kernel

# System Boot

---

- Bootstrap program (loader) locates the kernel, loads it and starts the kernel.
  - This can be a two-step procedure.
  - Bootstrap program loads another more complex boot program
  - That boot program loads the kernel
- Then control is given to kernel.
- Kernel starts the environment and makes the computer ready to interact with the user (via a GUI or command shell).
- Details depend on the system